

Rust: memory safety and systems programming

Manos Pitsidianakis

Software Engineering, NTUA, 2017-18

Rust is a systems programming language, like C, inspired by C and Ocaml.

Rust is somewhat Object Oriented.

- ▶ Traits instead of classes/interfaces
- ▶ Trait generics instead of inheritance

Types can implement Traits (what Rust calls interfaces+mixins).
We can abstract over many types by using Trait Objects (structs that implement a certain Trait).

```
pub trait Draw {  
    fn draw(&self);  
}
```

```
pub struct Screen {  
    pub components: Vec<Box<Draw>>,  
}
```

```
impl Screen {  
    pub fn run(&self) {  
        for component in self.components.iter() {  
            component.draw();  
        }  
    }  
}
```

Figure 1: A simple trait definition

```
pub struct Screen<T: Draw> {  
    pub components: Vec<T>,  
}  
  
impl<T> Screen<T>  
    where T: Draw {  
    pub fn run(&self) {  
        for component in self.components.iter() {  
            component.draw();  
        }  
    }  
}
```

Figure 2: Alternative definition using trait generics

Rust is safe unlike C, and also unsafe like C. This is because Rust is memory safe by default, but we can turn off memory safety and use it (somewhat) like C.

Undefined behaviour is also very limited.

Examples of memory safety:

- ▶ buffer overflow
- ▶ null pointer dereference
- ▶ use after free
- ▶ use of uninitialized memory
- ▶ illegal free

Unsafe Rust is code wrapped in an unsafe {} block. Its main difference from safe Rust is the ability to dereference raw pointers and calling external functions.

Wrapping unsafe code in special sections of our code enables easier checking for program correctness; only unsafe blocks must be checked for memory safety. This makes review very easy.

Safe Rust guarantees memory safety at compile-time.
A compiled safe Rust program will not segfault (unless there's an implementation error in the language itself).

The concept of lifetimes is the most interesting and confusing concept about Rust

- ▶ All values have one owner (reference) at a time.
- ▶ When the owner goes out of scope, the value's destructor is called.

To transfer values between functions, a value can either be

- ▶ moved, or
- ▶ borrowed as a reference (a kind of smart pointer)

```

fn foo() {
    let mut data = vec!['a', 'b', 'c']; // --+ 'scope
    capitalize(&mut data[..]);           //  |
//  ^~~~~~ 'lifetime //  |
    data.push('d');                      //  |
    data.push('e');                      //  |
    data.push('f');                      //  |
} // <-----+

fn capitalize(data: &mut [char]) {
    // do something
}

```

Figure 3:

To ensure memory safety at any time the following holds: For each value there's one mutable reference XOR unlimited immutable references.

Borrow checker checks these constraints during compilation. It's the first and hardest obstacle for newcomers to Rust. Not even unsafe Rust can turn off the borrow checker.

```
fn take(v : Vec<i32>) {  
    // ownership of the vector now transferred  
    // to v in this scope  
}  
  
let v = vec![1, 2, 3]; // take ownership  
take(v);               // moved ownership  
println!("v[0] is {}", v[0]);  
// error: use of moved value: `v`  
// println!("v[0] is {}", v[0]);  
//                               ^
```

Figure 4:


```
// note we're taking a reference, &Vec<i32>
// instead of Vec<i32>
fn take(v : &Vec<i32>) {
    // no need to deallocate the vector after
    // we go out of scope here
}

let v = vec![1, 2, 3]; // take ownership
// notice we're passing a reference, &v,
// instead of v
take(&v); // borrow ownership
// so now we have the ownership back
// we can use v again!
println!("v[0] is {}", v[0]);
// v[0] is 1
```

```
fn take(v : &Vec<i32>) {  
    v.push(5);  
}  
  
let v = vec![];  
take(&v);  
// cannot borrow immutable borrowed content `*v` as mutable  
// v.push(5);  
// ^
```

Figure 6:

```
// v is a mutable reference
fn take(v : &mut Vec<i32>) {
    v.push(5);
}

// v has to be mutable too in order for you
// to borrow a mutable ref
let mut v = vec![];
take(&mut v);
println!("v[0] is {}", v[0]);
// v[0] is 5
```

Figure 7:

Fearless concurrency.

“Do not communicate by sharing memory; instead, share memory by communicating.”

Message passing

```
// Suppose chan: Channel<Vec<i32>>  
  
let mut vec = Vec::new();  
// do some computation  
send(&chan, vec);  
print_vec(&vec);  
//Error: use of moved value `vec`
```

Figure 8:

Locks - Shared memory

Necessary evil for systems programming

The problem is not sharing state, but sharing it accidentally.

Reminder: Borrow checker doesn't allow more than one mutable reference at any time.

Data can be placed inside mutexes and shared between threads

Locking a mutex returns a mutable reference to the data inside.

Still only one mutable reference at any time.

What Rust doesn't prevent

memory leak by repeatedly allocating objects in an infinite loop

What Rust doesn't prevent

deadlocks

What Rust doesn't prevent

integer overflows (but it can panic or wrap around if we want)

What Rust doesn't prevent

logic errors, etc

References

<https://doc.rust-lang.org/book/>