

Hands-on Introduction to Bootloaders

Software Engineering, NTUA, 9th semester

Nikos Gavalas

March 2018

Sections

- Booting Process Overview
- Rolling our own bootloader
- BIOS Interrupts
- 32-bit Mode
- Global Descriptor Table
- Interacting with h/w
- Beyond the bootloader

Booting Process Overview

What happens when we press the power button on our computer? - 4 basic stages

1. BIOS/UEFI
2. Bootloader
3. Kernel
4. Init Process

BIOS (Basic Input/Output System)

- **Firmware** that performs basic diagnostic checks and initialization of the hardware during the booting process
- Searches for the “boot block”
 - MBR Partitioning → Looks for the “**Master Boot Record**” - The first 512 bytes of the first sector of the disk
 - The MBR contains the code that loads the bootloader afterwards.
- Limitations of the MBR:
 - Can't handle disks with more than 2 TB of space
 - Only supports up to four primary partitions

UEFI (Unified Extensible Firmware Interface)

- Successor of BIOS
- Most modern computers use it
- GPT disk partitioning
 - However GPT is backwards compatible, it's first sector is reserved for MBR code → booting is possible from computers with no UEFI.

Bootloader

- Responsible for passing parameters and **loading** the OS Kernel
- Famous Bootloaders: GRUB, U-Boot...
- The Kernel can't "see" the hardware initially (missing drivers)
 - Solution → mounts a **root filesystem** which contains all the needed modules to communicate with the hardware
 - this root fs was "initrd", now "initramfs"
 - the kernel creates a "root device" and mounts the "root partition" in read-only mode, then verifies the integrity of the other partitions with "fsck", and remounts the root partition in rw mode.
 - finally executes the "**Init**" process (systemd → initializes OS services) - every other process is a fork of "init".

Rolling our own bootloader

We'll be using the following tools:

- **nasm** to assemble our code
- **qemu** to emulate a pc (arch: 32-bit i386 - x86)

and as bootable storage media for qemu we'll emulate a **floppy disk**→ (way) simpler

Rolling our own bootloader

So when we power up our computer:

- BIOS is loaded from some flash memory on the motherboard
- Performs h/w checks
- Tries to find bootable storage devices (according to a boot sequence)
- Checks their first 512 bytes stored (MBR)
- If those bytes are **bootable**, it jumps to memory address `0x7c00` and starts executing, that's where we'll be placing the code for our bootloader

But what code is “bootable”?

- If the last 2 bytes of the MBR data are `0xaa55`, then the code is considered bootable

Rolling our own bootloader

Notice the address - 0x7c00 → 16-bit

It's because the processor at this point runs at **16-bit mode** ([Real Mode](#))

- Only 16-bit registers are available (ax, bx, ...)
- 20-bit segmented memory address space → 1 MiB of addressable memory
- No memory protection, access to all addressable memory

Later we will see how to enter the 32-bit “Protected Mode”

Getting our bootloader to print a message

Wait a second... print *where*, *how*?

We will use the [BIOS interrupt calls](#), which provide basic I/O functionality.

We can see that we can call BIOS code to print a character (`al ← char`) in tty mode, using the interrupt vector 0x10 (and `ah ← 0x0e`)

But first, a quick refresher on asm x86

All x86 registers [here](#)

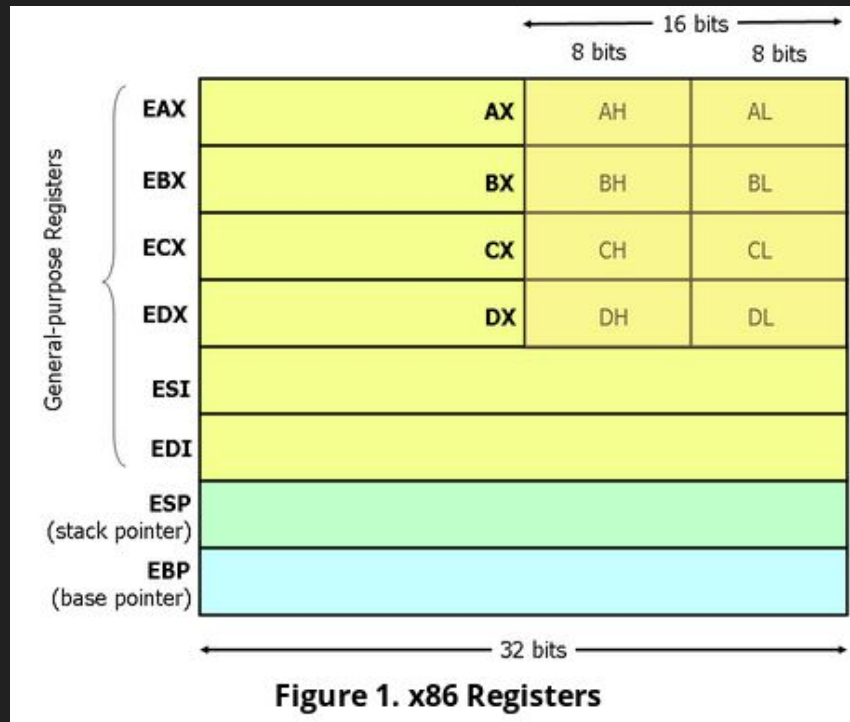
Control Registers

CRO

bit	label	description
0	pe	protected mode enable
1	mp	monitor co-processor
2	em	emulation
3	ts	task switched
4	et	extension type
5	ne	numeric error
16	wp	write protect
18	am	alignment mask
29	nw	not-write through
30	cd	cache disable
31	pg	paging

Index Registers

32 bit	16 bit	description
esi	si	source index
edi	di	destination index

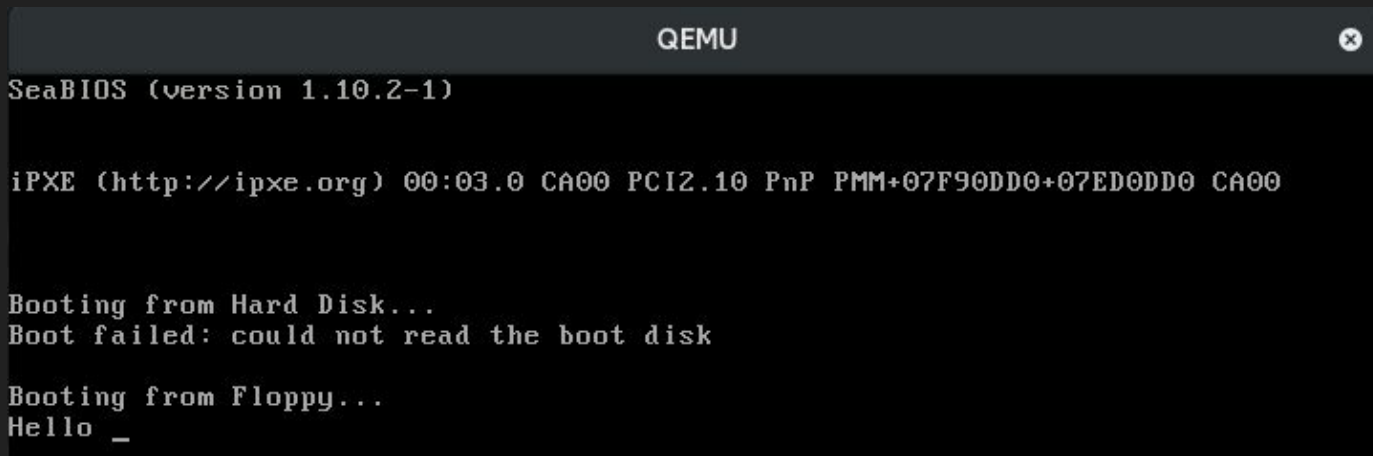


First steps

```
1  bits    16
2  org     0x7c00
3
4  global  main
5
6  main:
7      mov     bx, str           ; load str address
8
9  loop:
10     mov     byte ax, [bx]      ; load byte from where bx points to
11
12     inc     bx
13
14     push    bx
15     mov     ah, 0x0e           ; print char in TTY mode
16     int     0x10              ; s/w intr -> BIOS video services
17     pop     bx
18
19     cmp     al, 0              ; if we reached null
20     jz      halt
21     jmp     loop
22
23  halt:
24     cli                     ; clear interrupts
25     hlt                     ; halt
26
27  str:
28     db      "Hello", 0
29
30     times   510 - ($-$$) db 0 ; pad remaining 510 bytes with zeroes
31                                     ; (in nasm, $ -> current address,
32                                     ; and $$ -> beginning of current section)
33     dw      0xaa55            ; mbr magic number (2 bytes + 510 = 512 total)
```

Assemble and run

- Assemble with nasm:
 - `$ nasm bootloader.asm`
- Run with qemu:
 - `$ qemu-system-i386 -fda bootloader`

A screenshot of a QEMU window titled "QEMU" with a close button in the top right corner. The window displays a text-based boot process. It starts with "SeaBIOS (version 1.10.2-1)", followed by "iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+07F90DD0+07ED0DD0 CA00". Then it says "Booting from Hard Disk..." and "Boot failed: could not read the boot disk". Next, it says "Booting from Floppy..." and "Hello _".

```
QEMU
SeaBIOS (version 1.10.2-1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+07F90DD0+07ED0DD0 CA00

Booting from Hard Disk...
Boot failed: could not read the boot disk

Booting from Floppy...
Hello _
```

Let's check the assembled 512-byte binary

bootloader																																		
00000000	BB	14	7C	8B	07	43	53	B4	0E	CD	10	5B	3C	00	74	02	EB	F1	FA	F4	48	65	6C	6C	6F	00	00	00	00	00	00	00CS....[<.t....Hello.....	
00000020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
00000040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
00000060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
00000080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
000000a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
000000c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
000000e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
00000100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
00000120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
00000140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
00000160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
00000180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
000001a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
000001c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
000001e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	55	AAU.		
0000200																																		

We can see the magic number at the last two bytes

Something interesting

Let's check out our own computer's MBR code

- `$ dd if=/dev/sda bs=512 count=1 > mbr`

We can even disassemble it with objdump and take a closer look:

- `$ objdump -D -b binary -mi386 -Maddr16,data16 mbr`

Accessing more memory and 32-bit protected mode

To be able to access more memory we'll need to enable the [A20 line](#)

We can do that with another BIOS interrupt call → [int 0x15](#)

```
mov    ax, 0x2401
int     0x15      ; enable A20 bit
```

Now, to enter the Protected Mode and access the full 32-bit registers, we first need to set up a GDT ([Global Descriptor Table](#))

- The GDT contains entries telling the CPU about memory segments
- Useful for memory protection

GDT

Define memory segments:

See [here](#).

```
22 gdt_start:
23     dq  0x0
24
25 gdt_code:                                ; code segment from 0 to 0xffff
26     dw  0xFFFF
27     dw  0x0
28     db  0x0
29     db  10011010b                        ; r/w/e
30     db  11001111b
31     db  0x0
32
33 gdt_data:                                ; data segment from 0 to 0xffff
34     dw  0xFFFF
35     dw  0x0
36     db  0x0
37     db  10010010b                        ; r/w
38     db  11001111b
39     db  0x0
40
41 gdt_end:
42
43 gdt_pointer:                            ; pointer structure needed from the cpu to know
44                                         ; how big the gdt is
45     dw  gdt_end - gdt_start
46     dd  gdt_start
47
48
49 CODE_SEG equ gdt_code - gdt_start
50 DATA_SEG equ gdt_data - gdt_start
```

32-bit mode

Special command to load the GDT:

```
lgdt [gdt_pointer]
```

Finally we can enable protected mode by setting the corresponding bit on the special cr0 register:

```
mov eax, cr0  
or  eax, 0x1  
mov cr0, eax
```

We can now use the full registers

Writing on the VGA buffer

```
63      mov     bx, str           ; load str address
64
65      mov     esi, str
66      mov     ebx, 0xb8000      ; vga buffer address
67  loop:
68      lodsb                    ; load byte at address DS:(E)SI into AL
69      or      al, al           ; update flags
70      jz      halt
71      or      eax, 0x0200      ; green color
72      mov     word [ebx], ax    ; write to the vga buffer
73      add     ebx, 2           ; go to the next cell
74      jmp     loop
75
76  halt:
77      cli                        ; clear interrupts
78      hlt                      ; halt
79
80  str:
81      db     "Hello", 0
```

It's cleaner than before, isn't it



How about we access more hardware?

Let's go for the disk now

We don't really have a choice, because BIOS only loads the first 512 bytes, therefore if we want to write bigger programs, we 'll have to load them from the disk

How? You guessed it, [BIOS Disk Interrupts](#)

With int 0x13, ah=0x02, we can read sectors from the drive:

```
13      mov     [disk], dl      ; BIOS places the disk index on dl on startup, so we save it
14
15      mov     ah, 0x2         ; read sectors
16      mov     al, 1           ; sectors to read
17      mov     ch, 0           ; cylinder index
18      mov     dh, 0           ; head index
19      mov     cl, 2           ; sector index
20      mov     dl, [disk]      ; disk index
21      mov     bx, copy_target ; target pointer
22      int     0x13
23
```

What's next?

We can go further and start building a basic kernel

But how? We need a way to compile code (mainly C) for our new system

What we need is a *cross-compiling toolchain* → to compile software for our targeted arch from our host system

A toolchain includes everything that cross-compilation needs:

- Binutils (assembler, linker, ...)
- C compiler, C lib and more

However...

Building a cross-compiling toolchain is a nightmare, too difficult task

We use pre-compiled ones... but they don't offer much flexibility

The best choice is to go for a toolchain building tool, which makes the process of building a toolchain easier

The most popular toolchain building tools are [crosstool-ng](#), OpenEmbedded, Buildroot etc

To sum up

... There are a lot of ways to print 'Hello' on the screen...

Let's go through the steps once more

- BIOS searches for bootable code
- This bootable code in the boot sector loads the bootloader
- The bootloader loads the kernel
- The kernel checks the disks partitions, then remounts the filesystems, and executes init

Resources

- <https://deltahacker.gr/how-pcs-boot/>
- <https://wiki.osdev.org/Bootloader>
- https://wiki.osdev.org/Boot_Sequence
- https://wiki.osdev.org/Rolling_Your_Own_Bootloader
- <https://www.codeproject.com/Articles/664165/Writing-a-boot-loader-in-Assembly-and-C-Part>
- <https://osandamalith.com/2015/10/26/writing-a-bootloader/>

Questions?